

Maximilian Witte

# **Image Data for Machine Learning Research in Marketing**

EMAC Odense 2023



Universität Hamburg

DER FORSCHUNG | DER LEHRE | DER BILDUNG

# Why Use Image Data with Machine Learning for Marketing Research?



- 1 More than 6.5 billion images shared every day<sup>1</sup>
- 2 More than 60% of US adults use instagram<sup>2</sup>



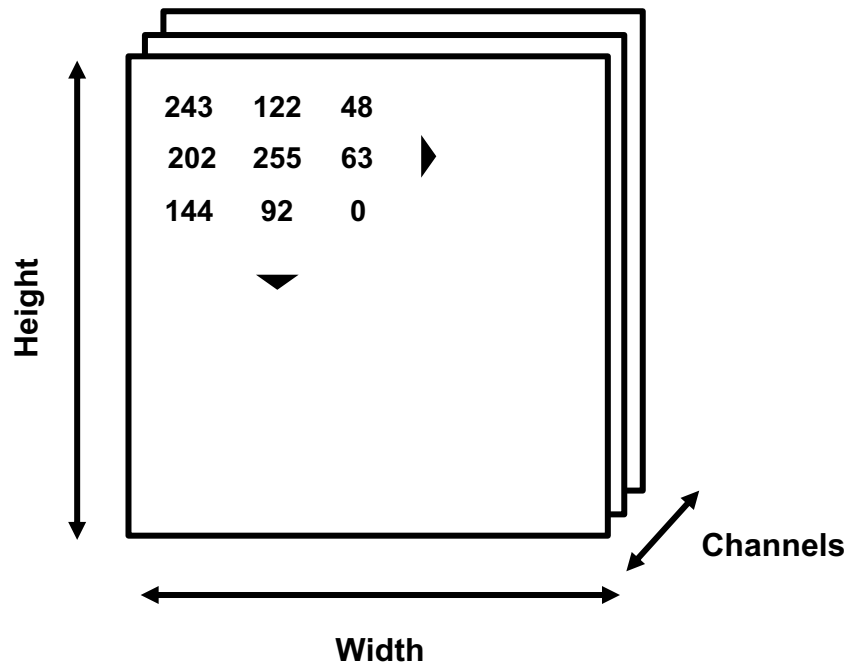
Automatically extract information for analysis

Predict behavior

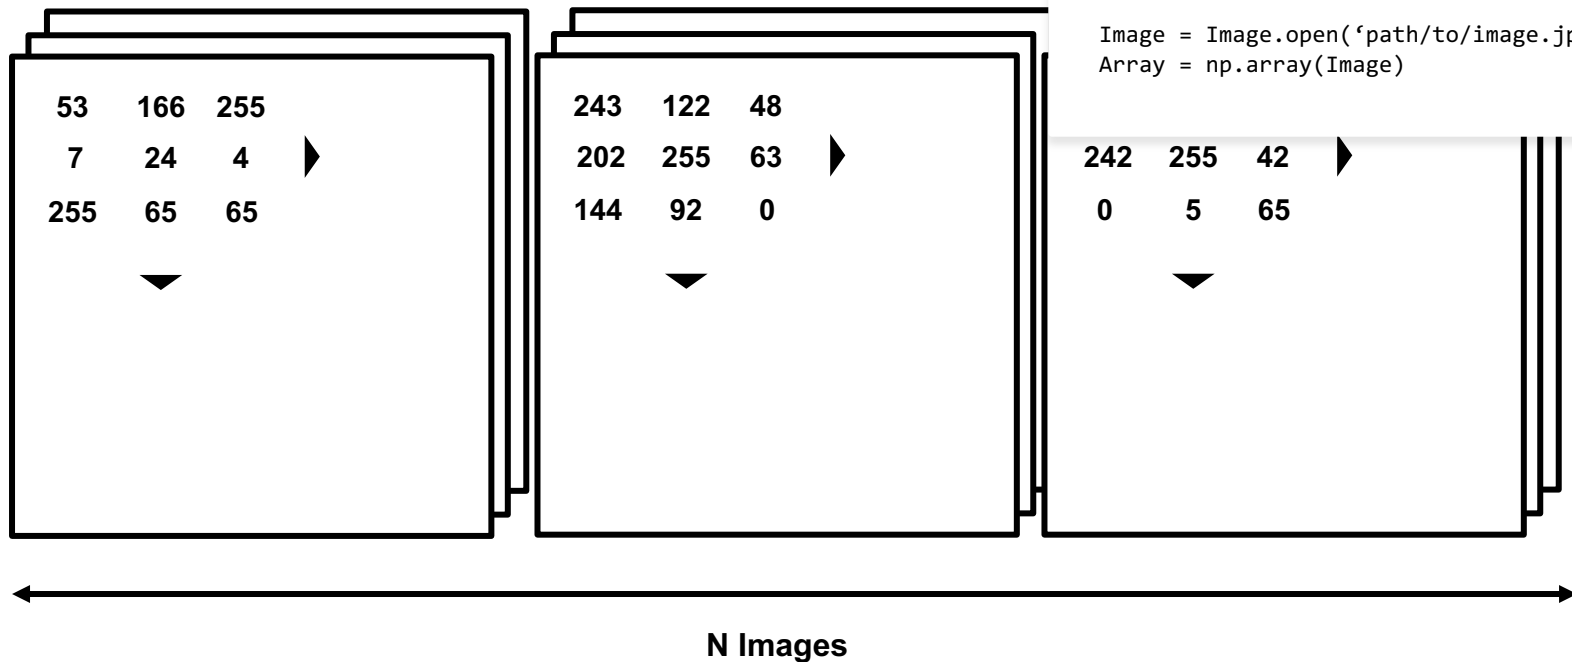


Goal always **generalization** on unseen data

# The 4 dimensions of image data



# The 4 dimensions of image data



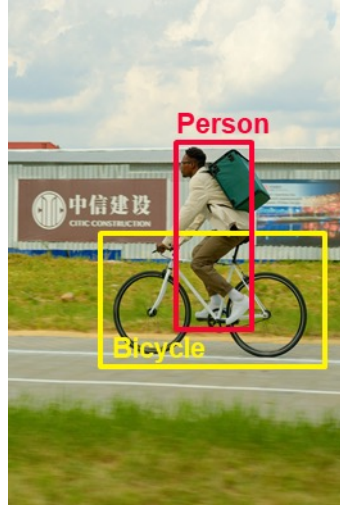
```
from PIL import Image  
import numpy as np
```

```
Image = Image.open('path/to/image.jpg')  
Array = np.array(Image)
```

# A multitude of potential tasks for image data



**Classification**



**Object  
Detection**



**Segmentation**

# A multitude of potential tasks for image data


**A**

Image Analysis

Classification



Object  
Detection



Segmentation



+

Depth  
estimation

Image-to-  
Text

...

**B**


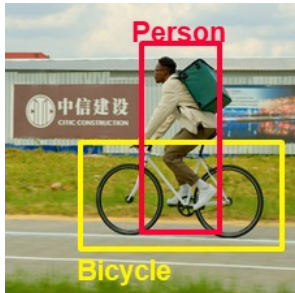

Generative AI<sup>1</sup>

Text2Image



+ Img2Img

# Tools to annotate your data

Task	Classification, Generative AI	Object detection	Segmentation, depth detection
			
Format	Folders, Tables	JSON, XML	Image, JSON
Tools	<ol style="list-style-type: none"> <li>1. Move images to folder</li> <li>2. PhotoSift<sup>1</sup></li> <li>3. Use Python GUI</li> </ol>	<ol style="list-style-type: none"> <li>1. Labelme<sup>2</sup></li> <li>2. V7<sup>3</sup></li> <li>3. Labelbox<sup>4</sup></li> </ol>	<ol style="list-style-type: none"> <li>1. Labelme<sup>2</sup></li> <li>2. V7<sup>3</sup></li> <li>3. Labelbox<sup>4</sup></li> </ol>

Python labeling  
notebook



[https://github.com/Maximilianwte/Image-labeling-tutorial/blob/main/Image\\_labeling.ipynb](https://github.com/Maximilianwte/Image-labeling-tutorial/blob/main/Image_labeling.ipynb)

1 PhotoSift (freeware) available for windows: <https://www.rvision.com/photosift/about.php>  
 2 Labelme (freeware) available for windows, mac and linux: <https://github.com/wkentaro/labelme>  
 3 V7 (free for education) available online: <https://www.v7labs.com>  
 4 Labelbox (freemium) available online: <https://labelbox.com/>



# Using Images for analysis

Table 3. Summary Statistics of Field Data.

Variable	Twitter		Instagram		Display Ads	
	Mean	SD	Mean	SD	Mean	SD
Dependent Variables						
Likes	5.06	233.79	255.33	4,380.30	—	—
Comments	.43	7.00	9.85	65.53	—	—
Purchase Intentions	.02	.15	.51	2.72	—	—
Brand-Image Type						
Brand Selfie <sup>a</sup>	28.73	—	19.63	—	72.68	—
Consumer Selfie <sup>a</sup>	7.93	—	11.57	—	18.80	—
Packshot <sup>a</sup>	63.35	—	68.80	—	8.51	—
Image Characteristics						
Logo Size	.12	.14	.07	.10	.33	.21
Logo Centrality	.63	.23	.63	.19	.81	.14
Visual Complexity	-1.96	.43	-2.09	.55	-1.40	.94
Brightness	.51	.15	.56	.13	.60	.13
Brightness Contrast	.25	.05	.26	.05	.24	.05
Post Characteristics						
Number of Words	2.20	.55	3.24	.82	2.29	.22
Number of Hashtags	.39	.59	2.29	.92	—	—
Number of Handletags	.23	.39	.25	.47	—	—
Branded Caption <sup>a</sup>	6.23	—	12.10	—	—	—
Branded Tag <sup>a</sup>	13.02	—	61.65	—	—	—
Ad Tag <sup>a</sup>	.35	—	2.35	—	—	—
First-Person Pronoun <sup>a</sup>	35.66	—	28.97	—	.13	—
Second-Person Pronoun <sup>a</sup>	14.59	—	16.79	—	35.25	—
Question Word Share	.01	.04	.01	.03	.01	.03
Netspeak Word Share	.02	.06	.01	.03	.00	.01
Positive Sentiment <sup>a</sup>	42.43	—	50.84	—	6.12	—
Neutral Sentiment <sup>a</sup>	46.25	—	45.24	—	92.73	—
Negative Sentiment <sup>a</sup>	11.32	—	3.92	—	1.15	—
Post Age	5.04	.64	2.39	1.48	—	—

Source: Hartmann, J., Heitmann, M., Schamp, C., & Netzer, O. (2021). The power of brand selfies. *Journal of Marketing Research*, 58(6), 1159-1177.



# Choosing a library for deep learning



**Developed by  
Google**

**Keras High-level  
API**

**Support for GPU's  
and TPU's**



**Developed by Meta**

**Pythonic syntax**

**Steep learning  
curve**



**HUGGING FACE**

**Large amount of  
pre-trained models**

**Easy-to-use**

**Support for PyTorch  
and TensorFlow**

# Preparing the data

The goal of ML is generalization →  
Simulate situation of having unseen data

## 1 Load data

## 2 Data Splitting

- **Train-test-split:** Split off percentage of training data for testing
- **K-fold cross-validation:** Split data in k equal datasets. Iterate over all subsets using each as validation dataset
- **Bootstrapping:** Create artificial datasets by randomly sampling observations

## 3 Feature extraction



```
from datasets import load_dataset

dataset = load_dataset("imagefolder",
                        data_dir="path/to/directory")
```



```
from sklearn.model_selection import
train_test_split

train_dataset, test_dataset =
train_test_split(dataset, test_size=0.2,
                  random_state=42)
```



```
from transformers import ViTFeatureExtractor

feature_extractor = ViTFeatureExtractor.
from_pretrained('google/vit-base')

def prepare(example_batch):
    inputs = feature_extractor([x for x in
                                example_batch['image']],
                                return_tensors='pt')
    inputs['labels'] = example_batch['label']
    return inputs

prepared_ds = dataset.with_transform(prepare)
```

# Increasing data variability by using Image Augmentations

Use augmentations to efficiently  
increase variability of the training  
data



More variability in data leads to  
better generalization



```
from torchvision import transforms

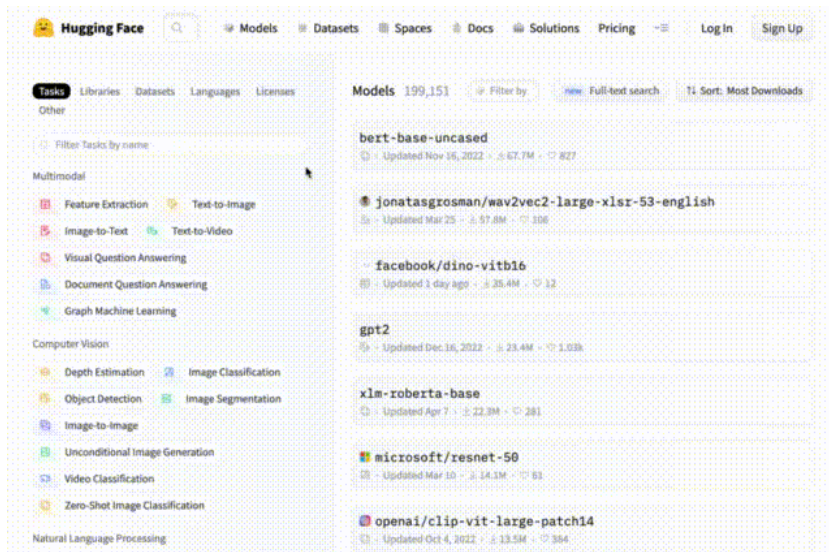
transform = transforms.Compose([
    transforms.ColorJitter(hue=0.5),
    transforms.RandomRotation(30),
    transforms.RandomHorizontalFlip(),
    transforms.Resize((512, 512))
])

img_transformed = transform(img)
```

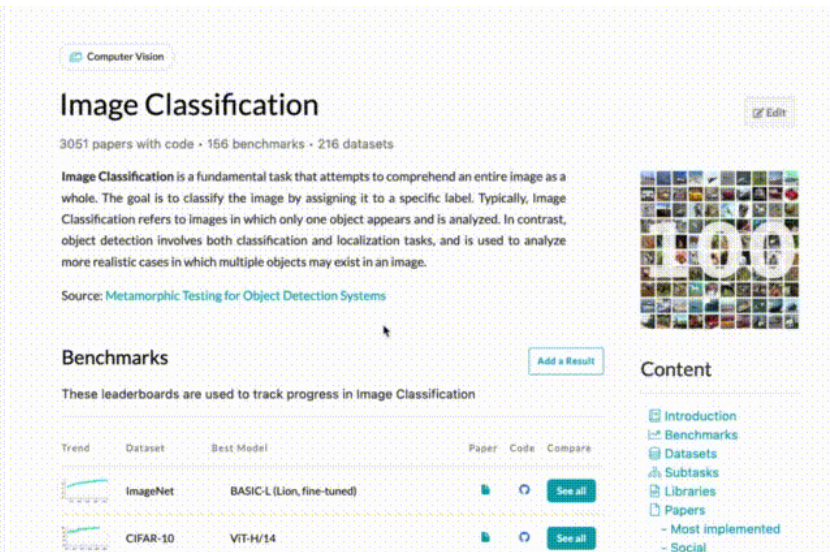
# Resources to find a good model



 Papers With Code



The screenshot shows the Hugging Face website's 'Models' section. The left sidebar contains navigation links for 'Tasks' (Feature Extraction, Image-to-Text, Visual Question Answering, Document Question Answering, Graph Machine Learning) and 'Computer Vision' (Depth Estimation, Image Classification, Object Detection, Image Segmentation, Image-to-image, Unconditional Image Generation, Video Classification, Zero-Shot Image Classification). The main content area lists several models, including 'bert-base-uncased', 'jonatasgrosman/wav2vec2-large-xlsr-53-english', 'facebook/dino-vitb16', 'gpt2', 'xlm-roberta-base', 'microsoft/resnet-50', and 'openai/clip-vit-large-patch14'. Each model entry includes its name, a brief description, and a 'Filter by' button.



The screenshot shows the 'Image Classification' page on the Papers With Code website. The page title is 'Image Classification' with a subtitle '3051 papers with code · 156 benchmarks · 216 datasets'. The main text describes the task: 'Image Classification is a fundamental task that attempts to comprehend an entire image as a whole. The goal is to classify the image by assigning it to a specific label. Typically, Image Classification refers to images in which only one object appears and is analyzed. In contrast, object detection involves both classification and localization tasks, and is used to analyze more realistic cases in which multiple objects may exist in an image.' Below this, there is a 'Benchmarks' section with a table showing leaderboards for ImageNet and CIFAR-10. The table has columns for 'Trend', 'Dataset', 'Best Model', 'Paper', 'Code', and 'Compare'. The 'ImageNet' row shows 'BASIC-L (Lion, fine-tuned)' as the best model, and the 'CIFAR-10' row shows 'ViT-H/14'. To the right of the main text is a 'Content' sidebar with links to 'Introduction', 'Benchmarks', 'Datasets', 'Subtasks', 'Libraries', 'Papers', 'Most implemented', and 'Social'.

# Training the model

- Process of training the model
- Question of optimal hyperparameters for training



- Grid search, Random search, Bayesian Optimization
- Automated libraries for tuning hyperparameters

## Training (without automated Hyperparameter tuning)



```
from transformers import
ViTForImageClassification

model =
    ViTForImageClassification.from_pretrain
    ned("google/vit-base").to('cuda')

training_args = TrainingArguments(
    per_device_train_batch_size=64,
    evaluation_strategy="steps",
    num_train_epochs=20,
    logging_steps=50,
    learning_rate=1e-4,
)

trainer = Trainer(
    model=model,
    args=training_args)

train_results = trainer.train()
```

## Training (with automated Hyperparameter tuning)



```
import optuna
from transformers import
ViTForImageClassification

model =
    ViTForImageClassification.
    ned("google/vit-base").to('cuda')

def objective(trial):
    training_args = TrainingArguments(
        per_device_train_batch_size=trial.suggest_c
        ategorical('batch_size', [16,32,64]),
        evaluation_strategy="steps",
        num_train_epochs=20,
        logging_steps=50,
        learning_rate=trial.suggest_float('lr', 1e-
        5, 1e-3),
        load_best_model_at_end=True,
    )
    trainer = Trainer(
        model=model,
        args=training_args)
    train_results = trainer.train()
    return train_results['eval_accuracy']

study =
    optuna.create_study(direction='maximize')
study.optimize(objective, n_trials=100)
study.best_params
```

Python image  
classification



[https://github.com/Maximilianwte/Image-tutorial/blob/main/Image\\_Classification.ipynb](https://github.com/Maximilianwte/Image-tutorial/blob/main/Image_Classification.ipynb)


# Finally using your model to predict on new data

Model



Inside

3.75



Outside

4.52



```
from transformers import AutoImageProcessor
from transformers import ViTForImageClassification
from PIL import Image

image_processor = AutoImageProcessor.from_pretrained("google/vit-base")

model = ViTForImageClassification.from_pretrained("/path/to/finetuned/model")

image = Image.open("/path/to/image")
inputs = image_processor(image, return_tensors="pt")
with torch.no_grad():
    logits = model(**inputs).logits

predicted_label = logits.argmax(-1).item()
print('Label: ' + model.config.id2label[predicted_label])
```

Label: 'Car'


Python image classification



[https://github.com/Maximilianwte/image-classification-tutorial/blob/main/image\\_classification.py](https://github.com/Maximilianwte/image-classification-tutorial/blob/main/image_classification.py)

# B

## Using Generative AI



# Diffusion models revolutionizing generative computer vision since 2022

Closed Source & Paid



DALL-E 2

<http://labs.openai.com/auth/login>



Midjourney

<https://discord.com/invite/midjourney>

Open Source & Free



Stable Diffusion



Python



DeepFloyd



Python




# Generating images can be done from text prompts or with an input image to transform

Prompting



Img2Img



Inpainting



Outpainting



# Generating Images using Stable Diffusion



```
!pip install diffusers
import torch
from diffusers import StableDiffusionPipeline

pipe = StableDiffusionPipeline.from_pretrained("CompVis/stable-diffusion-v1-5",
        revision="fp16", torch_dtype=torch.float16, use_auth_token=True).to("cuda")

images = pipe(prompt="A dog writing a novel, cartoon", height=512, width=512,
        guidance_scale=7.5).images
```

images



Stable diffusion  
notebook



[https://github.com/Maximilianwte/Image-tutorial/blob/main/Stable\\_Diffusion.ipynb](https://github.com/Maximilianwte/Image-tutorial/blob/main/Stable_Diffusion.ipynb)

# Getting your data into the model

## A Adapting an Input Image



Possible with all models

## B Fine-tuning the weights and embeddings of the model



```
pipe(prompt="Photo of xyzperson holding a  
handcreme", height=512, width=512,  
guidance_scale=7.5).images[0]
```



Currently only with Stable Diffusion  
and DeepFloyd

Fine-tuning SD  
Dreambooth



[https://github.com/MaximilianWitte/Image-tutorial/blob/main/Finetuning\\_Stable-Diffusion\\_Dreambooth.ipynb](https://github.com/MaximilianWitte/Image-tutorial/blob/main/Finetuning_Stable-Diffusion_Dreambooth.ipynb)

# A few principles on working with machine learning models

**1** Generalization is key

**2** Leverage pre-trained models

**3** Better data trumps a larger model

**4** Get fitting architecture and loss function

**5** Prefer the simpler model

# Resources

## Presentation GitHub Repository



<https://github.com/Maximilianwte/Image-tutorial>

## Blogs




<https://huggingface.co/docs>



<https://towardsdatascience.com/>

## Book Recommendation



<https://www.manning.com/books/deep-learning-with-python>

# Contact



**Maximilian Witte**

Chair of Marketing &  
Customer Insight

University of Hamburg

**#Unstructured data**

**#Generative AI**

**#Decision Aids**



[maximilian.witte@uni-hamburg.de](mailto:maximilian.witte@uni-hamburg.de)



[linkedin.com/in/maximilianwte](https://linkedin.com/in/maximilianwte)

